



(a) sequential: GPU idle during host work, governor holds 306 MHz



(b) prefetch: decode of $i+1$ overlaps GPU of i ; clock rises (911 MHz)



(c) full: two decoders, rows of $i-1$ overlap GPU of i ; GPU near 1020 MHz

